

Protecting Public Data From AI Using Dataset Poisoning

Josie Barber

Division of Science and Mathematics
University of Minnesota, Morris

November 12, 2025

Greg Rutkowski



[@greg_rutkowski on Instagram 2024]



[Rutkowski, 2021]

Outline

Background

Overview of Models

Training Data

Types of Attacks

Attack Techniques and Case Studies

Ethical Considerations


Conclusion

Background Overview

- ▶ Generative AI systems (LLMs and diffusion models) are vulnerable to manipulation at multiple stages.
- ▶ Small manipulations can produce large behavioral changes.
- ▶ Understanding these vulnerabilities requires knowledge of:
 - ▶ Model structures and lifecycles
 - ▶ Attack types across training, inference, and deployment phases

Generative AI Model Life-Cycle

- ▶ **Training phase:** Dataset compilation, preprocessing, and parameter optimization.
- ▶ **Inference phase:** Model applied to new inputs; fine-tuning or prompt engineering may occur.
- ▶ **Deployment phase:** Integration into real-world systems; potential monitoring and retraining.



Types of Generative AI Models

Diffusion Models


- ▶ **Used for:**
 - ▶ Upscaling images
 - ▶ Text-to-image/video generation
 - ▶ Style transfer and animation
 - ▶ Stable Diffusion
 - ▶ Midjourney

Large Language Models (LLMs)

- ▶ **Used for:**
 - ▶ Text & Code generation
 - ▶ Questioning & Reasoning
 - ▶ Translation and summarization
 - ▶ OpenAI GPT
 - ▶ Google Gemini

Diffusion Models


- ▶ Learn to reverse a gradual noise process to reconstruct realistic data.
- ▶ Starts by gradually corrupting training data until it is nothing but noise



[Higham et al. 2025]

Diffusion Models

- ▶ Iteratively denoise inputs: each step slightly refines the output.
- ▶ Score the resulting data and test it against original images



[Higham et al. 2025]

Large Language Models (LLMs)

- ▶ Learns to predict the next token in a sequence, modeling patterns and relationships in text.
- ▶ Designed to mainly “pay attention” to which words in a sentence relate to each other.
- ▶ Capture meaning from context rather than fixed word definitions, allowing them to generalize across topics.

Training Data

- ▶ Sources of Training Data:
 - ▶ Books and academic articles
 - ▶ Encyclopedias
 - ▶ Code repositories
 - ▶ User-generated content
 - ▶ Web-scraping
- ▶ Remove duplicates, low-quality, or otherwise harmful content
- ▶ Synthetic data or data augmentation


Ethics: Training Data

- ▶ Bias in Data
- ▶ Privacy Concerns
- ▶ Harmful Content
- ▶ Transparency
- ▶ Representation Gaps
- ▶ Misinformation
- ▶ Outdated Content
- ▶ Copyright Issues

Copyright Issues and Transparency

Poisoning Attacks

- ▶ Intentional modifications during model training to change model output
 - ▶ Called an "adversarial attack"
- ▶ Common forms:
 - ▶ Label-flipping
 - ▶ Clean-label modification
 - ▶ Trojan/backdoor triggers
- ▶ Can embed hidden behaviors or induce biases



[Xu et al. 2025]

Offensive vs Defensive Techniques

Offensive


- ▶ Protecting data by adding malicious functions if dataset is used
 - ▶ Embedding triggers to entirely sabotage models
 - ▶ Inserting fake or misleading information to make generalization unusable

Defensive

- ▶ Changes to protect against the unwanted learning from a dataset more directly
 - ▶ Embedding triggers to add watermarks to output
 - ▶ Making Unlearnable Examples to block learning entirely



Unlearnable Training Data

- ▶ Human readable, however unreadable to a model
- ▶ Methods:
 - ▶ **Visual/format perturbations**
 - ▶ **Glyph substitutions**
 - ▶ **Adversarial noise**



Watermarking Public Datasets

- ▶ Embed verifiable signatures without affecting dataset utility.
- ▶ Tang et al. framework:
 - ▶ Clean-label
 - ▶ Adversarial perturbations
 - ▶ Backdoor triggers



Original Data

Perturbed Data

Colorful Patch

Texture Pattern

[Tang et al. 2025]

Ethical Considerations

- ▶ **Dual use:** Same methods lead to both good and bad
- ▶ **Transparency:** Poisoned data can be misleading or low quality
- ▶ **Accountability:** Poisoned data is difficult to trace or retract.
- ▶ **Consent:** Poisoned data is difficult to trace or retract.

Towards Responsible AI

- ▶ Encourage **informed consent**
- ▶ Use **transparent protection methods**
- ▶ Build **ethical standards**
- ▶ Balance **privacy defense**

Questions?

References

- ▶ Higham, Catherine F., et al. 2025. Diffusion Models for Generative Artificial Intelligence: An Introduction for Applied Mathematicians.
- ▶ Huang, et al. 2021. Unlearnable Examples: Making Personal Data Unexploitable.
- ▶ Rutkowski, Greg. 2024. Instagram.com.
- ▶ Rutkowski, Greg. Chaos Dragon. 2021. Art Station.
- ▶ Tang, Ruixiang, et al. 2025. Watermarking Public Datasets for Tracing Data Usage in Machine Learning.
- ▶ Xu, Wenrui, et al. 2025. A Survey of Attacks on Large Language Models.